

INDEPENDENT INSTITUTIONAL

Product Readiness Assessment

Auriva Healthcare Operating System

Commercial-launch audit | 23 July 2026 | Evidence reviewed: code, tests, product artefacts, release documents, and current competitor materials

MANDATE This is not an implementation sign-off. It is an external product, clinical-workflow, UX, business and engineering review. The assessment challenges internal readiness claims and treats missing evidence as a risk.

FINAL LAUNCH DECISION: ⚠ NEEDS SIGNIFICANT WORK

Audit scope, method and limits

This assessment is evidence-led. It reviewed 55 application pages, 110 API route handlers, 74 automated test files, the Prisma domain model, operating/debt/release documentation, and five internal Product Office workflow audits. A fresh test run completed successfully; TypeScript type-check completed successfully. The project lint command failed with 43 errors and 51 warnings. The production build could not complete in this restricted environment because Next font fetching failed; this is classified as a build reproducibility risk, not proof the product cannot build in a connected production environment.

The audit attempted to render the local application in the controlled browser. The browser could not reach the local development server, therefore visual findings are source- and component-led rather than a pixel-level rendered inspection. Scores are intentionally conservative where runtime evidence was unavailable.

Evidence hierarchy

- Runtime/automated check result
- Implemented code/schema/API route
- Product and release artefact
- Official competitor material
- Inference or strategic recommendation

Evidence sources

- Internal: docs/AURIVA-PRODUCT-BIBLE.md; docs/AURIVA-RELEASE-1X-ROADMAP.md; docs/RELEASE-CANDIDATE.md; docs/technical-debt.md; docs/product-office-audit/*; docs/knowledge-base/*; src/app, src/components, src/services and prisma/schema.prisma.
- External official competitor sources current at audit date: Jane feature/booking pages; Cliniko online payments; SimplePractice client portal and billing automation; Halaxy calendar; Practo Ray; DocPulse; Clinicea features; Healthie pricing. URLs are listed in Appendix A.

1. Executive Summary

BOARD CONCLUSION Auriva is not launch-ready as a commercial Healthcare Operating System. It is a serious, well-architected, early product with an unusually coherent encounter-to-cash thesis; it is not yet a dependable, purchasable clinic operating system across the patient, revenue, clinical and launch-support lifecycle.

The most valuable thing Auriva has is not its breadth. It is the decision to make a clinic day move through one shared appointment state machine: booking/walk-in → waiting → consultation → checkout. That is a credible operational core. The problem is that almost every purchase-critical loop around it is incomplete: patient activation depends on unproven production OTP; the product does not communicate outside the app; patients cannot pay online; intake and consent are absent; scheduling is shallow; analytics are thin; clinical data is largely unstructured; pricing is not finalized; and release gates remain open.

Dimension	Score	Status	Board finding
Overall Product Score	5.2	RED	A credible encounter-to-cash foundation, not a commercially complete Healthcare OS.
Market Readiness	4.0	RED	Category-table stakes—messaging, online payments, intake/consent, reporting and mature scheduling—remain absent or thin.
UI	6.8	AMBER	Tokenized and intentional, but the assessment is source-led because the live UI was not reachable from the audit browser.
UX	5.4	AMBER	Happy-path stage ownership is strong; cross-workflow recovery, discovery and automation are incomplete.
Engineering Maturity	5.8	AMBER	Strong boundaries, audit/event seams and 74 test files; quality gates currently fail and scaling assumptions are explicit.
Business Readiness	3.7	RED	No finalized pricing, package design, proof of ROI, migration plan or support operating model.
Launch Readiness	3.5	RED	Production OTP, backups/restore, monitoring and release gates remain unchecked; lint fails.
Commercial Readiness	3.8	RED	Value proposition has promise but lacks a sellable outcome package and hard switching proof.
Customer Delight	4.7	RED	The warm patient shell is promising, but it does not communicate, collect payment, or guide care outside the clinic.
Clinical Readiness	4.8	RED	Useful documentation is present, yet structured clinical governance, consent, prescribing safeguards and longitudinal record depth are immature.

Two-page executive conclusion

The internal materials repeatedly use “mature,” “production-ready” and “release candidate.” Those terms are defensible only for a constrained single-clinic, staff-assisted encounter-to-cash pilot after infrastructure gates are closed. They are not defensible for a broad commercial launch, a patient-led product, a multi-clinic platform or an enterprise edition. The report therefore rejects a launch narrative that outruns the finished experience.

Auriva should not respond by adding every competitor feature. That would violate its own guardrails and lose its best strategic asset: operational simplicity. It should instead complete the smallest number of end-to-end loops that make a selected clinic buyer visibly better off every day: reliable identity/communications, pre-visit readiness, front-desk exception management, structured clinically safe documentation, checkout/payment, follow-up/next booking, and owner proof of value.

The five non-negotiable launch blockers

- Production OTP delivery, delivery observability and patient-login proof.
- Backups, restore drill, environment validation, monitoring/alerting and operational ownership.
- A clean quality gate: lint, type check, tests and production build must be reproducible.
- Consent-based outbound communication and at least one patient-reachable reminder/document pathway.
- A sellable commercial package: chosen beachhead, finalized price/limits, assisted onboarding and pilot success definition.

Final verdict

⚠ Needs Significant Work. Auriva is appropriate for tightly controlled design-partner pilots only after P0 work is closed. It should not be called “Launch Ready,” should not sell Enterprise, and should not promise a complete Healthcare OS in market messaging.

2. Product Vision Review

The vision is unusually clear in prose: an India-first independent clinic should run a calm room, capture cash reliably and make patients feel cared for. The six-pillar boundary is disciplined, and the explicit refusal to become HRMS/ERP is strategically correct. The positioning, however, is too large for the delivered product. “Healthcare Operating System” invites comparison with products that already offer mature communications, online payment, forms/consent, reports, recurring scheduling, telehealth and broader clinical record functionality.

Question	Board answer	Required correction
Is the vision clear?	Yes. The room-flow and revenue-leak thesis is memorable.	Preserve it; turn it into an observable promise.
Is positioning obvious?	Partly. The category is clear, the buyer and beachhead are not.	Name a clinic type, specialty and daily outcome.
Would customers understand it?	They understand appointments/queue/billing; “OS” can confuse.	Lead with a concrete before/after workflow.
Would investors understand it?	They will see ambition but ask why incumbents do not win.	Show segment evidence, retention mechanism and unit economics.
Would doctors understand it?	They will understand the workbench but reject empty promises.	Demonstrate safety, speed and post-visit continuity.

3. Target Customer Review

Segment	Fit now	Board judgment	Why
Solo GP/consulting doctor	Amber	Pilot only	Core appointment/consult/billing works, but patient engagement and discovery are thin.
Physiotherapy	Amber → Green after C1/C2	Best beachhead candidate	Multi-session plans and follow-up leakage map directly to roadmap.
Dentistry	Amber → Green after C1/C2	Strong but not now	Treatment plans, procedures, consent and imaging/doc needs are not ready.
Psychology	Red	Avoid initially	Needs intake, consent, secure messaging, telehealth and specialty documentation.
Nutritionists	Red	Avoid initially	Healthie/Jane strength is care plans, messaging and remote engagement; Auriva is not there.
Single-site multi-doctor clinic	Amber	Controlled pilot	Best fit for queue/reception thesis once P0s close.
Polyclinic/chain	Red	Do not sell	Rooms/resources, deep multi-clinic, enterprise controls and reporting are not mature.

4. Market Readiness

Auriva’s potential differentiation is operationally coherent, not feature breadth: a shared appointment state machine, India-first phone/family/UPI assumptions and a deliberately warm patient experience. It is not differentiated enough today to overcome competitor table stakes. Jane connects online booking, reminders, intake and payments; Cliniko can collect online payment and issue payment links; SimplePractice automates billing document delivery and portal

payments; Halaxy connects online bookings, reminders, records, invoicing and reporting; Practo Ray promotes SMS reminders, online payment and treatment plans; Clinicea advertises automated consent, reminders, packages, telehealth and patient access. Auriva must win a narrower workflow, not claim parity.

Competitor	What it establishes as table stakes	Auriva gap / response
Cliniko	Mobile-friendly online payments, payment links and automated invoices.	No online payment; prioritize India-appropriate UPI/payment flow.
Jane	Online booking, reminders, waitlist, intake/consent, payments, secure portal and scheduling depth.	Do not compete broadly; match pre-visit + communication basics for chosen segment.
Healthie	Telehealth, scheduling, business tools and wellness/team workflows.	Avoid remote-care/nutrition market until care plans and messaging mature.
Halaxy	Online booking, email/SMS reminders, patient records, invoices and financial reports.	Communication and owner reporting are non-negotiable gaps.
SimplePractice	Portal, online payment, billing automation, documentation/search/filter depth.	Patient self-service and information retrieval are thin.
Practo Ray	India-specific reminders, online payment, EMR/treatment plans and mobile access.	Auriva must prove lower friction/front-desk superiority, not generic parity.
DocPulse	Broad integrated HMS/clinic suite incl. EMR, billing, labs, pharmacy, reminders, analytics.	Do not chase hospital breadth; own independent clinic flow.
Clinicea	Consent, plans/packages, telehealth, communication, multi-location and deep specialty tooling.	Auriva needs a beachhead and simpler proof, not a feature race.

5. Complete User Journey Audit

Stage	What works	Break/failure point	Recommendation
Account/clinic creation	Organization/workspace foundation is real.	No proven first-value, import or guided operational configuration.	Instrument setup and offer assisted pilot onboarding.
Login/identity	Staff sessions, forced password change and OTP design are thoughtful.	Production OTP delivery is not complete; patient beta is blocked.	Close R-001 before exposing patient journey.
Booking	Slot booking and server conflict checks exist.	No forms, consent, deposit, reminder or waitlist recovery.	Complete pre-visit loop.
Walk-in/reception	Fast minimal capture and status machine are strong.	Search/disambiguation, reassign and exception recovery are weak.	Surface Health ID/DOB; add transfer.
Queue/consult	Clear state ownership minimizes ambiguity.	Direct consult bypass and shallow escalation can create operational drift.	Guard exceptions and display recovery.
Clinical documentation	Shared workbench and factual safety chips are promising.	Free-text encounter data, no signed amendment model or clinical governance.	Structure core data and govern releases.
Prescription/labs	Print and basic order/result paths exist.	No full medication safety, clinical review or integration depth.	Avoid over-claiming clinical maturity.
Checkout/payment	Invoice/payment/correction architecture is a real strength.	No online payment, tax active fields or unified service-to-charge proof.	Close payment and service capture loops.
Records/documents	Timeline/document platform foundation is good.	Search, delivery, patient-friendly access and longitudinal completeness are weak.	Ship timeline + communication together.
Notifications/follow-up	In-app notification exists.	No reachable notifications; no automated follow-up.	This is a launch blocker, not polish.
Settings/logout	Role/session handling is robust in intent.	Several surfaces expose “Soon,” confusing real vs planned actions.	Hide/defer incomplete controls.

6. Role-Based Journey Review

Role	Ease / cognitive load	Breakpoints	Board recommendation
Doctor	Good queue-to-consult focus; workbench can be compelling.	Search weak; clinical structure/safety immature; placeholders harm trust.	Make Workbench one validated, high-confidence clinical surface.
Receptionist	Best-designed role conceptually; stage ownership is intuitive.	Search, reassignment, calendar context, communications and payment recovery are incomplete.	Make reception the beachhead and test it under real traffic.
Clinic owner	Has command-center/team/billing foundation.	Cannot yet see a measurable business case or mature reports.	Sell outcomes dashboard, not administration.
Patient	Warm framing and family model are promising.	Login, communication, payments and after-visit support are weak.	Build active, not passive, patient relationship.
Administrator	Strong RBAC/audit/event scaffolding.	Some admin UI is internal/platform-like; enterprise controls incomplete.	Keep admin scope narrow until enterprise needs are real.

7. UI Review

The implemented design system is more mature than the business product. It has semantic colors, light/dark tokens, common components, focus-visible rules and reduced-motion handling. The intended Pine/Honey separation can be a distinct brand asset. The weakness is product truth: polished cards and intentional copy cannot compensate for disabled

“Soon” navigation, informational-only controls, and operational journeys that stop before patient communication/payment. A premium product feels complete in the moments people pay for, not merely consistent in its component code.

8. UX Review

Mental model strength: the appointment state machine gives each role a visible handoff. This is excellent and should remain sacrosanct. Mental model failure: user intent is repeatedly acknowledged but not completed—pause booking, insurance, upcoming capabilities, communications, advanced schedule management. Each is a small expectation breach. The UX roadmap should prioritize recovery and closure over decorative refinement.

9. Healthcare Workflow Review

A real receptionist can plausibly process a walk-in, move a patient through the queue and collect payment. That is a meaningful vertical slice. A real clinic, however, also needs certainty before arrival, handling for late/absent/changed doctors, room/capacity decisions, patient reminders, rebooking, consent, follow-up and owner visibility. The product handles the center of a clinic day better than its edges. Commercially, the edges are where no-shows, liability and churn happen.

10. Business Model Review

There is an edition shape (Solo / Professional / Enterprise) but no approved pricing. That is not pricing readiness. Feature gates alone are not packaging. A viable near-term commercial offer is: an assisted, single-clinic Professional pilot for one validated specialty, with a fixed implementation fee, named service limits, core staff seats, committed support and a 90-day outcome review. Do not offer Enterprise until shared rate limits, clinic-scoped grants, multi-clinic analytics, security/compliance evidence, support SLAs and proven customer operations exist.

Revenue lever	Readiness	Board view
Subscription tiers	Red	Shape exists, commercial terms do not.
Implementation/onboarding	Red	No mature migration or success operating model.
Communications usage	Red	High potential after consented delivery exists.
Payment processing	Red	High potential but no patient payment capability.
Treatment plans/packages	Amber	Strong future revenue lever; needs full workflow first.
Multi-clinic/enterprise	Red	Do not monetize prematurely.
Analytics/reporting	Amber	Renewal lever once decision-grade metrics exist.

11. Quality Assurance Audit

Positive evidence: 74 test files, 704 test declarations, 110 API routes, centralized status transitions, audit/event seams, health/readiness endpoints and documented operational debt. Negative evidence matters more for launch: the current lint run fails with 43 errors/51 warnings; UI-level E2E proof is not evidenced; production build reproducibility is sensitive to remote font fetching in this environment; single-instance rate limiting/reminder sweeps and unpaged list APIs are documented limitations. The test count is not a release decision. Gate on scenario coverage, severity and clean commands.

12. Accessibility Audit

Positive signals include focus-visible styling, several ARIA labels, semantic status labels, reduced-motion CSS, shared controls and dark-mode tokens. Gaps: no automated accessibility audit evidence, no keyboard workflow certification, no screen-reader transcript evidence, no touch-target audit and no contrast evidence across all dynamic states. Accessibility is promising in implementation intent, not verified in launch assurance.

13. Design System Audit

The system has genuine reuse: shared controls, semantic CSS variables, dark mode, state components and a clear staff/patient tone. This is a comparative strength. It needs a governed token inventory, visual regression tests, icon/status usage rules and fewer product-level exceptions. The design system should protect the operational core, not create a false sense of maturity around unfinished workflows.

14. Product Maturity

Classification	Board assessment
MVP	Yes: a real staff-assisted encounter-to-cash slice exists.

Classification	Board assessment
Friendly-clinic beta	Conditional: only after P0 launch/infrastructure, quality and communication gates close.
Launch candidate	No: commercial, patient and operational completeness is insufficient.
Production ready	Only narrow internal/single-clinic components, not the commercial product.
Enterprise ready	No: do not market or sell it.

15. Future Roadmap Review

The stated sequence—Treatment Planning → Procedure Management → Clinical Timeline → Prescription Platform → Communication—is strategically sound but incomplete as a launch plan. Communication must move ahead of broad patient/pilot claims because it is the delivery mechanism for the existing promise. Treatment planning is the best commercial growth step only if the beachhead is physio/dental; otherwise it is premature specialization. Do not add generic “revenue platform” breadth. Make services-to-invoice, payment, rebooking and outcome reporting reliable first.

Roadmap item	Board verdict	Required change
Revenue platform	Keep narrow	Complete payment, tax decision, service capture, correction and owner reporting—not accounting.
Treatment planning	Prioritize for selected specialty	Require plan-to-session-to-payment-to-next-booking end-to-end metric.
Procedure management	Do after plans	Avoid specialty hardcoding; validate session model with design partners.
Clinical timeline	Accelerate	Make it doctor pre-consult and patient record truth, with performance/search.
Prescription platform	Accelerate with governance	Structured Rx, signing/amendment and safety checks—not AI diagnosis.
Communication platform	Move to launch-critical	Consent, templates, delivery, preference and audit is required before public beta.

16. Launch Readiness

LAUNCH BLOCKED The product should not launch broadly. The release documentation itself lists unchecked production database, backup/restore, production OTP/SMS, production configuration and monitoring/alerting gates. This audit adds a failing lint gate, absent commercial packaging, no outbound patient communication, no online patient payment and insufficient E2E/clinical governance evidence.

17. Commercial Readiness

Would clinics pay? A narrow group might pay for an assisted pilot if Auriva proves fewer missed handoffs, faster walk-ins and fewer uncollected visits. Would they switch? Not yet at scale: data migration, communication, payment and recurring scheduling are table stakes. Would users recommend it? Receptionists could recommend the queue board; patients will not recommend a portal that does not reliably reach them, accept payment or guide them after the visit. Stickiness will come from longitudinal care plans, documents/communication, a trusted patient record and owner ROI—not from the admin architecture.

18. SWOT Analysis

Strengths	Weaknesses	Opportunities	Threats
One status machine across roles; clear product guardrails; strong billing/audit foundation; warm patient design intent.	Unfinished patient communication/payment/pre-visit loop; weak commercial readiness; shallow scheduling/search/analytics; clinical data debt.	Own India-first single-site clinic flow; win multi-visit specialty; prove outcomes; turn communications/timeline into retention.	Feature-complete incumbents; scope dilution into ERP/HIS; clinical/privacy incident; no reference customers; expectation debt.

19. Risk Register

Risk class	Status	Risk	Consequence	Owner	Mitigation
Business	Red	No beachhead/pricing/ROI	Cannot sell repeatably	CPO	R-008/R-028/R-064
Product	Red	Patient loop stops at portal	Low adoption/no-show value unrealized	VP Product	R-005/R-009/R-043
UX	Amber	Front-desk exception paths missing	Stress and workaround behavior	UX Director	R-013/R-019/R-039
Technical	Red	Quality gates fail; single-instance dependencies	Regression/reliability risk	CTO	R-003/R-006/R-007
Clinical	Red	Unstructured records, no governance/consent	Safety and liability risk	Clinical Safety Board	R-004/R-015/R-017
Commercial	Red	Enterprise claims ahead of capability	Churn/reputation risk	GTM Lead	R-030/R-098
Security/privacy	Amber	Compliance pack and scale hardening incomplete	Procurement/security risk	Security Lead	R-020/R-054/R-056

20. Top 100 Recommendations

The register is sequenced by launch impact. “Critical” means a launch/pilot gate; “High” means the first commercial roadmap; “Medium” means required to protect adoption and scale; “Low” means valuable only after the core is proven. Complexity: S ≤1 sprint, M 1–2, L 2–4, XL multi-milestone.

Critical — P0

Critical recommendation detail 1 of 2

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-001	Launch	Live OTP delivery and delivery observability are an unchecked external-beta gate.	Patients cannot reliably log in; every patient-facing promise collapses.	Integrate one production SMS/OTP vendor, delivery callbacks, retries, rate limits and a daily failure dashboard.	Makes patient activation measurable and protects trust.	M / P0
R-002	Launch	Backups and a production restore rehearsal are not evidenced as complete.	A recoverability failure with health data is existential.	Provision encrypted backups, define RPO/RTO, execute a documented restore drill, and make it a release gate.	Prevents catastrophic loss and reassures institutional buyers.	M / P0
R-003	Engineering	The current lint command fails with 43 errors and 51 warnings.	A failing quality gate normalizes unsafe change and undermines release claims.	Fix all source lint errors; exclude archival prototypes deliberately; block merge and release on clean lint, types and tests.	Cuts regression cost and restores credibility.	M / P0
R-004	Clinical safety	No documented clinical governance, sign-off, or specialty validation process is visible.	A clinically unsafe workflow can create patient harm and liability.	Create a Clinical Safety Board, risk classification, change-control, usability validation and release sign-off for clinical changes.	Reduces clinical and legal risk; becomes an enterprise sales asset.	M / P0
R-005	Patient engagement	No outbound SMS, WhatsApp, email or push delivery exists; in-app alerts require patients to open Auriva.	No-show reduction, follow-up completion and document delivery cannot happen.	Ship consented multi-channel communications with templates, preferences, delivery status and cost controls before broad launch.	Unlocks retention, attendance and visible daily value.	L / P0

Critical recommendation detail 2 of 2

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-006	QA	There is no evidenced browser end-to-end release suite for the full care day.	Service tests do not prove the real clinic can move a patient through the UI.	Automate role-based golden paths: setup, booking, walk-in, queue, consult, invoice, payment, portal and permission denial.	Reduces launch-day workflow breaks.	L / P0
R-007	Security	In-memory rate limiting and reminder sweep are single-instance only.	Scale can weaken abuse protection and duplicate operational actions.	Move rate limits, idempotency and scheduled delivery locks to shared infrastructure before multi-instance deployment.	Prevents security and reliability incidents during growth.	M / P0
R-008	Commercial	Pricing is explicitly not finalized.	Sales cannot package, quote or forecast a product without a defensible commercial offer.	Define editions, limits, add-ons, implementation fees, usage-metered messaging and a price-testing plan.	Enables selling, forecasting and controlled expansion.	M / P0
R-009	Patient experience	Patients cannot pay online or from the portal.	The financial journey stops at reception; self-service and collection efficiency are lost.	Launch payment links/UPI intent/QR with receipts, reconciliation and failure handling.	Accelerates collection and makes the portal useful.	L / P0
R-010	Findability	Reception and doctor surfaces lack a global patient search; lists have no pagination.	Busy staff lose time and risk opening the wrong patient as data grows.	Implement tenant-scoped global search by name, phone, health ID and DOB; add paging, server filtering and audit context.	Cuts handling time and improves safety.	L / P0

High — P1

High recommendation detail 1 of 4

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-011	Onboarding	Clinic setup lacks proven time-to-first-value instrumentation and guided data migration.	A buyer may never reach a successful first day.	Create an activation checklist, sample data, import tools, onboarding telemetry and concierge playbook.	Improves conversion and lowers support cost.	M / P1
R-012	Scheduling	Calendar is day-only on reception; advanced recurrence, drag/drop and scheduling requests are deferred.	Auriva loses to commodity scheduling tools before its operations story is experienced.	Prioritize recurring schedules, reschedule, conflict explanations, cancellation policy and waitlist before broad rollout.	Reduces staff workload and no-shows.	L / P1
R-013	Reception	Doctor assignment is fixed at creation; the board has no reassignment or queue transfer.	One late doctor creates avoidable queue congestion and poor patient experience.	Add a guarded reassign/transfer action with reason, patient notification and audit record.	Improves throughput without adding headcount.	M / P1
R-014	Revenue	Consult services/treatments are not fully wired into a reliable consult-to-invoice commercial flow.	Revenue capture is shallow and dependent on manual work.	Make ordered services the canonical invoice input; show expected charge, waiver and exception reasons at handoff.	Raises capture rate and billing confidence.	L / P1
R-015	Clinical	Clinical documentation remains largely free text on the Appointment “god table.”	Longitudinal care, safe analytics and interoperability are constrained.	Define a versioned clinical encounter model with structured diagnoses, vitals, problems, medications and	Creates an enduring clinical data asset.	XL / P1

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
				provenance.		

High recommendation detail 2 of 4

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-016	Prescription	Prescription safety is incomplete: no interaction checks, structured medication governance or pharmacy workflow.	Medication errors and support burden remain material.	Ship structured Rx documents, allergy/duplicate therapy checks, immutable signing, amendment trail and local formulary strategy.	Raises clinical trust and differentiates doctor workflow.	L / P1
R-017	Consent	Digital intake, consent and signature workflows are absent.	Clinical and medico-legal documentation starts too late and remains weak.	Introduce configurable forms, versioned consent, e-signature, completion status and pre-visit gating.	Cuts front-desk work and improves defensibility.	L / P1
R-018	Analytics	Command Center is an activity view, not a decision system; analytics are explicitly thin.	Owners cannot see whether Auriva improves cash, flow or retention.	Ship a small KPI pack: no-show, waiting time, consult completion, unpaid balance, rebooking and plan conversion.	Creates owner stickiness and renewal evidence.	M / P1
R-019	Operations	Queue wait aging is visual only; capacity thresholds and intervention workflows are deferred.	The product displays congestion without helping resolve it.	Set configurable service-level thresholds, escalation states and one-action recovery options.	Protects patient experience and clinic throughput.	M / P1
R-020	Data protection	No formal compliance posture, privacy mapping or data-processing pack is evidenced.	Institutional buyers will block procurement; trust claims are unsubstantiated.	Produce data-flow map, retention/deletion policy, breach response, vendor register, DPA and jurisdiction-specific legal review.	Unlocks procurement and reduces legal exposure.	M / P1

High recommendation detail 3 of 4

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-021	Product	Treatment planning and procedure tracking are roadmap items, not current value.	Physio and dental customers cannot run their core multi-visit care model.	Complete plan → session → checkout → next-booking loop before targeting these specialties.	Opens high-LTV specialty segments.	L / P1
R-022	Product	Clinical timeline is not yet the decisive pre-consult context screen.	Doctors still assemble history across surfaces.	Unify visits, documents, plans, labs, payments and active problems with role-specific views.	Reduces consult prep time.	L / P1
R-023	UX	“Soon” and reserved features appear in several live role surfaces.	They signal incompleteness at the exact moment a buyer evaluates value.	Hide non-sellable navigation by entitlement; replace placeholders with meaningful next actions.	Improves perceived quality and task focus.	S / P1
R-024	Support	No mature customer-success operating model or implementation SLA is evidenced.	First clinics will experience adoption risk as product failure.	Define pilot success plan, escalation paths, response targets, weekly value review and feedback intake.	Improves retention and referenceability.	M / P1
R-025	Migration	No legacy data import and reconciliation journey is visible.	Switching cost remains too high for established clinics.	Provide patient, appointment, balance and document imports with dry run, mapping and exception report.	Removes a primary purchase barrier.	L / P1

High recommendation detail 4 of 4

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-026	Payments	No documented tax/GST-ready commercial receipt strategy is evidenced; tax fields are inactive.	Indian clinics may not be able to use billing as their system of record.	Validate target segments; implement tax configuration, compliant invoice fields and accountant export where required.	Protects billing adoption.	M / P1
R-027	Roles	Nurse and technician permissions exist but dedicated action surfaces are deferred.	Role promise is broader than usable work; teams revert to shared credentials or paper.	Build minimum safe vitals capture and results-entry workflows or remove roles from packaging.	Improves team adoption and access safety.	M / P1
R-028	Market	Target segment is too broad: solo doctors through polyclinics, across many specialties.	Message, workflow and sales motion become generic.	Choose one beachhead (initially cash/UPI, multi-visit physio or dental) and say no to the rest.	Concentrates learning and GTM efficiency.	S / P1
R-029	Market	Mental-health and nutrition workflows have no explicit fit validation.	Competing against SimplePractice/Healt hie without forms, messaging or telehealth is unwinnable.	Avoid these segments until their must-have journey is built and validated.	Prevents expensive failed sales.	S / P1
R-030	Marketing	The product claims “Healthcare OS” before sufficient operational breadth is delivered.	This creates expectation debt and a credibility gap in demos.	Sell “front-desk-to-cash operating system for [beachhead] clinics” until the OS claims are earned.	Raises conversion and trust.	S / P1

Medium — P2

Medium recommendation detail 1 of 9

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-031	Accessibility	Accessibility is implemented through scattered ARIA/focus patterns, with no automated audit evidence.	Keyboard and assistive-technology failures will surface late.	Add axe-based CI, keyboard test scripts, focus-order acceptance criteria and screen-reader smoke tests.	Reduces exclusion and enterprise risk.	M / P2
R-032	Accessibility	Touch targets are often compact 28–36px controls in clinically busy areas.	Tablet and gloved/fast-use interaction errors increase.	Adopt a 44px minimum for primary touch actions and evaluate board density separately.	Fewer operational mistakes.	M / P2
R-033	Accessibility	Dark-mode semantic colors are declared but not contrast-verified across component states.	Clinical status may become ambiguous in night use.	Run WCAG contrast tests for all status foreground/background combinations and focus states.	Improves safety and compliance.	S / P2
R-034	UX	Error states are designed, but recovery ownership and retry semantics are inconsistent by workflow.	Users may retry unsafe mutations or abandon work.	Standardize mutation states: pending, success, failure, retry-safe, and support reference ID.	Reduces support and duplicate actions.	M / P2
R-035	UX	Autosave in clinical forms has no evidenced visible save confidence or conflict resolution.	Clinicians cannot know whether a critical note persisted.	Show save status, last saved time, unsaved-warning and concurrent-edit behavior.	Builds clinician trust.	M / P2

Medium recommendation detail 2 of 9

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-036	UX	The doctor “Pause booking” action is informational only.	A false control damages trust during a stressful day.	Remove it until real, or implement a scoped availability/booking hold with confirmation.	Eliminates expectation mismatch.	S / P2
R-037	Search	Health ID search exists in the backend but is not surfaced in core UI.	A designed safety identifier is not delivering operational value.	Expose Health ID and DOB-assisted disambiguation in reception and doctor search.	Reduces wrong-patient risk.	S / P2
R-038	Search	Billing and lab worklists lack patient/invoice/test search.	Staff must scan queues manually under pressure.	Add server-side search, active filters and persistent query state.	Cuts processing time.	M / P2
R-039	Scheduling	Reception board and calendar are split rather than one operational workspace.	Staff context-switches between “who is here” and “what is next.”	Provide a linked day workspace with board/capacity/calendar panes and preserved filter context.	Improves front-desk speed.	L / P2
R-040	Patient	Patient records cannot be searched.	Longitudinal records become unusable as history grows.	Add secure full-text or metadata search with document type/date filters.	Improves patient self-service.	L / P2

Medium recommendation detail 3 of 9

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-041	Patient	Family sharing lacks granular consent and proxy governance.	Family convenience may cause privacy breaches.	Add per-profile consent, guardian role, expiration and audit trail before marketing family care broadly.	Protects trust and compliance.	L / P2
R-042	Patient	The patient app shows medicines informationally with no adherence, refill or escalation loop.	The portal is passive after the visit.	Start with refill due and follow-up reminders, then validate adherence features by specialty.	Creates repeat engagement.	M / P2
R-043	Patient	There is no “after the visit” closing moment.	Patients leave without a clear plan, documents or next action.	Create a concise care summary with next appointment, Rx, payment status and help channel.	Raises trust and rebooking.	M / P2
R-044	Clinical	Lab result display is basic and “recent labs” is coming soon in doctor context.	Clinical information is incomplete at point of care.	Make results reviewable, trendable, attributable and acknowledged before claiming a lab workflow.	Improves safety and clinician value.	L / P2
R-045	Clinical	No specialty template strategy is evidenced beyond generic SOAP.	Generic documentation produces either low adoption or unsafe custom workarounds.	Develop validated specialty starter packs only for the selected beachhead, with governance.	Improves adoption and data quality.	M / P2

Medium recommendation detail 4 of 9

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-046	Clinical	No clear correction/amendment UX for signed clinical artifacts is evidenced.	Clinical records may be silently changed or hard to correct.	Adopt lock/sign/amend pattern with reason, timestamp, author and version view.	Protects medico-legal integrity.	M / P2
R-047	Clinical	No allergies/contraindications data-quality workflow is evidenced.	Safety chips can be stale or missing.	Add structured capture, source, last-reviewed date and high-risk acknowledgment.	Improves meaningful safety support.	M / P2
R-048	Design system	Semantic design tokens exist but brand color tokenization is documented as debt.	Visual changes risk inconsistent hard-coded implementation.	Finish semantic token inventory, component variants and visual regression baselines.	Speeds safe product evolution.	M / P2
R-049	Design system	Two header conventions are intentionally allowed but need usage rules.	Intentional variety can look accidental without explicit hierarchy.	Document page archetypes, use cases and review criteria; enforce in component examples.	Raises premium coherence.	S / P2
R-050	Design system	No observed visual regression test gate protects the claimed design system.	Small changes can erode hierarchy across 55 pages.	Capture approved states at desktop/tablet/mobile and diff in CI.	Protects brand and accessibility.	M / P2

Medium recommendation detail 5 of 9

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-051	Performance	Several list surfaces fetch full datasets and filter client-side.	Latency and browser load will degrade with real clinic history.	Set page limits, cursors, indexed queries and response budgets.	Supports scale without redesign.	L / P2
R-052	Performance	Command-center service has a documented N+1 per-clinic query pattern.	Multi-clinic owner experience will slow at precisely the upsell moment.	Replace with aggregate queries/read models and performance tests.	Supports enterprise conversion.	M / P2
R-053	Reliability	Demo seeding is non-transactional and not concurrency-locked.	Demos can fail visibly or misrepresent product reliability.	Use a durable demo tenant pipeline, lock and readiness check.	Improves sales demo confidence.	M / P2
R-054	Security	CSP nonce work is deferred.	The browser security posture is weaker than a healthcare product should accept long term.	Complete CSP hardening and include report-only rollout/violation monitoring.	Reduces XSS exposure.	M / P2
R-055	Security	Trusted proxy handling is configurable only as future debt.	IP-based protections can be bypassed/misapplied in production topology.	Document proxy topology and make trusted hops explicit, tested configuration.	Protects rate limiting and audit accuracy.	S / P2

Medium recommendation detail 6 of 9

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-056	Security	Some public-facing doctor directory reads create PII/exposure risk in the debt register.	Unnecessary data exposure harms patient/staff privacy.	Minimize fields, require intent, audit access and test enumeration resistance.	Reduces privacy risk.	M / P2
R-057	QA	Tests are heavily service/API oriented; usability acceptance criteria are thin.	A technically correct workflow can still be unusable at the front desk.	Add task-time, error-rate and completion criteria to release evidence.	Improves real-world quality.	M / P2
R-058	QA	No defect taxonomy tied to clinical severity is evidenced.	Teams may treat patient safety defects as ordinary UI bugs.	Define P0 clinical, P0 privacy, P1 revenue/flow and stop-ship rules.	Improves release decision quality.	S / P2
R-059	QA	Network/offline UI is claimed but no proof of safe offline mutations is evidenced.	Connectivity messaging without durable sync could create data loss expectation.	Specify offline scope; test queueing/conflicts or remove "sync when we can" claims.	Protects trust in low-connectivity settings.	L / P2
R-060	Architecture	Capability grants are clinic-agnostic.	Growing groups cannot safely tailor authority by branch.	Introduce clinic-scoped grants before enterprise packaging.	Makes multi-clinic authorization sellable.	L / P2

Medium recommendation detail 7 of 9

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-061	Architecture	Identity unification is deferred while phones are unique across account types.	Staff/patient role overlap creates future support and migration friction.	Design account/profile linkage migration now; avoid new phone-bound constraints.	Preserves platform flexibility.	L / P2
R-062	Architecture	Duplicate clinic/departments prevention is application-level only.	Concurrent writes can create duplicate master data.	Add database uniqueness constraints with data-cleanup migration.	Improves integrity.	M / P2
R-063	Architecture	The status machine allows scheduled → in consultation directly.	Reception handoff can be bypassed and operational truth becomes ambiguous.	Make bypass an explicit exception reason or remove it for standard flow.	Improves queue integrity.	S / P2
R-064	Commercial	No quantified ROI calculator or benchmark evidence is visible.	A clinic owner hears aspiration, not a buying case.	Instrument and publish baseline-to-outcome metrics from pilots.	Raises close rate and pricing power.	M / P2
R-065	Commercial	No clear free trial, assisted pilot or conversion motion is evidenced.	Acquisition funnel and support load cannot be planned.	Test a tightly scoped 30-day assisted pilot with success milestones and conversion trigger.	Creates repeatable GTM learning.	S / P2

Medium recommendation detail 8 of 9

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-066	Commercial	No partner/integration strategy is visible for payments, messaging, labs or accounting export.	Building all adjacent capabilities will dilute the product.	Publish integration principles and choose a small partner roadmap.	Accelerates completeness without ERP drift.	M / P2
R-067	GTM	The marketing site needs evidence-led claims, security posture and persona-specific proof.	Prospects cannot distinguish promise from implementation.	Create vertical landing pages, workflow demos, proof points, implementation expectations and transparent scope.	Improves qualified demand.	M / P2
R-068	GTM	No sales enablement for objection handling, comparison or implementation is visible.	Founder-led selling will not scale and claims will drift.	Create battlecards, demo script, security FAQ, migration FAQ and pricing guardrails.	Improves sales consistency.	S / P2
R-069	Operations	Room/slot capacity and resource scheduling are not operationally mature.	Polyclinic workflows will break even if appointments work.	Keep polyclinics out of target until rooms, resources and capacity policies are validated.	Avoids over-selling.	S / P2
R-070	Operations	No explicit cancellation/no-show recovery workflow is visible.	Lost capacity is not systematically recovered.	Add outcome reasons, automated rebooking prompt, waitlist offer and owner metric.	Recovers revenue.	M / P2

Medium recommendation detail 9 of 9

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-071	Operations	No explicit queue escalation for clinical urgency is visible.	Priority sorting without policy can be unsafe or arbitrary.	Define limited, auditable urgency flag roles and escalation policy with clinical governance.	Improves safety.	M / P2
R-072	Documents	Reserved document types are platform-ready but not content-ready.	Users cannot rely on a single document system for routine care.	Prioritize referral, certificate and lab request only after clinical/legal template review.	Improves completeness and stickiness.	M / P2
R-073	Documents	No document delivery/read receipt workflow exists.	Issued documents may not reach patients or create defensible evidence.	Link document sharing to communications, permissions, expiries and audit receipt.	Improves service and reduces calls.	M / P2
R-074	Reporting	No data-export/self-service report strategy is evidenced.	Owners and accountants will export manually or abandon the product.	Offer limited role-based CSV/PDF exports and scheduled operational reports.	Improves adoption without becoming ERP.	M / P2
R-075	Observability	Correlation-ID coverage remains incomplete.	Support investigations will be slow across real user flows.	Finish request correlation and standard operational dashboards/runbooks.	Cuts incident resolution time.	M / P2

Low — P3

Low recommendation detail 1 of 5

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-076	UX	No keyboard shortcut map is surfaced for high-frequency staff actions.	Experienced receptionists cannot reach speed potential.	Add discoverable shortcuts for search, walk-in, check-in, call-in and collect.	Raises throughput.	M / P3
R-077	UX	No saved views or user preferences for queues/worklists are visible.	Users repeatedly reconstruct the same context.	Allow role-safe saved filters and column density preferences.	Improves repeat productivity.	M / P3
R-078	UX	No comparative workload planning view is available to doctors.	Schedule pressure remains reactive.	Add simple day capacity and planned-vs-actual view after scheduling foundations are fixed.	Improves work-life and adoption.	M / P3
R-079	Patient	No patient feedback/review loop follows completed care.	The product misses a low-cost trust and growth signal.	Ask one short, consented experience question after a completed visit.	Creates learning and referral opportunity.	S / P3
R-080	Patient	Patient education content is absent.	A portal can feel transactional rather than caring.	Support clinic-authored pre/post-visit instructions only, not medical advice generation.	Raises perceived care quality.	M / P3

Low recommendation detail 2 of 5

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-081	Design	Iconography is sourced from a standard set but a documented usage guide is not visible.	Semantic variation can become inconsistent.	Publish icon meaning, status pairing and do-not-use rules.	Improves learnability.	S / P3
R-082	Design	Mobile responsive behavior is asserted but not independently evidenced by visual testing.	Desktop-first layouts may fail in patient contexts.	Create breakpoint acceptance snapshots and real-device pilot checks.	Reduces mobile churn.	M / P3
R-083	Engineering	External font fetching makes builds environment-sensitive.	Builds can fail in restricted networks, as this audit observed.	Self-host approved fonts or provide robust local fallback in the build pipeline.	Improves reproducibility.	S / P3
R-084	Engineering	The repository README remains the default Next.js starter text.	New engineers and evaluators lack an accurate operating entry point.	Replace with architecture, local setup, quality gates, test data and production runbook pointers.	Reduces onboarding cost.	S / P3
R-085	Engineering	Production readiness claims are dispersed across many documents.	Teams can miss the actual release truth.	Create one versioned release dashboard with evidence links and owner/date for each gate.	Improves governance.	S / P3

Low recommendation detail 3 of 5

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-086	Product	Doctor profile contains several "coming soon" tiles.	It dilutes professional credibility without advancing care.	Keep profile focused on bookable, verified attributes; defer aspirational features from live UI.	Improves trust.	S / P3
R-087	Product	Patient insurance settings expose a non-functional pathway.	Creates confusion in a market where coverage expectations vary.	Remove until supported, or clearly explain the no-insurance product policy.	Cuts support requests.	S / P3
R-088	Product	The public directory lacks specialty/location/availability filters.	Discovery does not serve patient intent at scale.	Add search ranking and filters once the target geography and provider supply model are defined.	Improves booking conversion.	M / P3
R-089	Product	No structured review-management workflow is visible despite doctor ratings.	Ratings can appear stale or unmanaged.	Define source, moderation, response and consent policy before using ratings commercially.	Protects brand trust.	M / P3
R-090	Product	Document categories are broad without a patient-friendly naming strategy.	Patients may not recognize clinical artifacts.	Use clear consumer labels, summaries and progressive disclosure in portal.	Improves comprehension.	S / P3

Low recommendation detail 4 of 5

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-091	Product	No ROI nudges at clinic milestones are visible.	Owners may not perceive the value created by operations improvement.	Show monthly avoided no-shows, recovered balances and rebooking outcomes when data is reliable.	Supports retention.	M / P3
R-092	Product	No in-product learning/help pattern is visible for first-time workflows.	“Zero training” is an unproven aspiration.	Use contextual, dismissible guidance and role-based first-run tasks.	Reduces onboarding friction.	M / P3
R-093	QA	No production synthetic monitoring of public booking/OTP is evidenced.	Failures may be discovered by patients first.	Run consent-safe synthetic checks and alert on booking/OTP degradation.	Improves reliability.	M / P3
R-094	QA	No load test thresholds for morning queue spikes are documented.	Performance risk is unmeasured at the highest-stress clinical moment.	Benchmark queue/search/checkout under representative concurrent load.	Protects clinic flow.	M / P3
R-095	QA	No accessibility statement or public conformance target is visible.	Enterprise buyers cannot assess inclusion maturity.	Publish supported standard, known limitations and feedback route after audit.	Improves procurement readiness.	S / P3

Low recommendation detail 5 of 5

ID	Area	Problem	Impact	Recommendation	Business value	Complexity / priority
R-096	Architecture	Event platform has no customer-visible consumer product yet.	Engineering sophistication has little commercial narrative.	Use events first for auditable communications and operational alerts, not generic platform messaging.	Converts foundation into value.	M / P3
R-097	Architecture	There is no explicit data archival strategy visible for high-volume clinical history.	Long-term performance and retention cost are unknown.	Define retention/archival policies, query boundaries and legal review.	Supports scale responsibly.	M / P3
R-098	Commercial	No formal enterprise SLA/support package is ready.	Enterprise tier risks being only a feature label.	Do not sell Enterprise until security review, support SLAs, governance and multi-clinic depth exist.	Avoids churn and reputational damage.	S / P3
R-099	GTM	No customer reference program is evident.	Auriva lacks proof beyond product narrative.	Recruit 3–5 design partners with explicit outcome measurement and reference rights.	Builds believable demand.	M / P3
R-100	Strategy	The roadmap sequence is sound in intent but lacks decision gates tied to customer evidence.	Teams can build elegant features without commercial pull.	Require a validated problem, target segment, success metric and stop condition before each milestone.	Improves capital allocation.	S / P3

21. Screen Review

Scores are source-led because the controlled browser could not reach the local development server. They assess implemented purpose, connected capability, known limitations and visible route/component intent—not a pixel-perfect live render.

Screen inventory 1 of 6

Screen	Route	Score	Strength	Weakness	Missing information / suggestion
Marketing home	/	6.0	Clear category story and visual system.	"Healthcare OS" promise exceeds present product breadth.	Use a concrete front-desk-to-cash proof story, pilot outcomes and scope boundaries.
About	/about	5.5	Brand mission can build empathy.	Mission is not buyer evidence.	Add team credibility, clinical governance and customer proof.
Platform	/platform	5.0	Pillar framing is coherent.	Claims need implementation-level substantiation.	Map each promise to shipped capability or remove it.
Solutions	/solutions	4.8	Possible segment orientation.	Broad segments hide poor beachhead focus.	Create one vertical page per validated segment only.
Industries	/industries	4.5	Signals ambition.	Cross-specialty claims exceed current specialized workflow depth.	Limit to validated specialties.
Pricing	/pricing	3.0	A pricing surface exists.	Commercial package is not finalized.	Do not expose until plans, limits and support are sellable.
Security	/security	5.0	Security concern is acknowledged.	Claims need a procurement-ready evidence pack.	Publish controls, scope and independent review status.
Trust	/trust	5.0	Trust is a useful brand theme.	Trust cannot substitute for communications, reliability and governance.	Use audited operational proof.
Compliance	/compliance	4.2	Signals seriousness.	Compliance posture is not evidenced as complete.	Obtain legal review and state exact jurisdictional position.
Customers	/customers	3.5	Correct destination for proof.	No demonstrated reference base is evident.	Populate only with attributable pilot outcomes.

Screen inventory 2 of 6

Screen	Route	Score	Strength	Weakness	Missing information / suggestion
Book demo	/book-demo	5.5	Appropriate enterprise CTA.	No qualification/expectation framework is shown.	Add target-fit questions and implementation promise.
Contact sales	/contact-sales	5.0	Clear route to human help.	Sales cannot answer pricing/migration confidence yet.	Equip sales before demand generation.
Get started	/get-started	4.5	Supports self-serve aspiration.	First-value and setup readiness are unproven.	Gate through assisted pilot until activation is proven.
Public booking	/book/[doctorId]	5.8	Real slot selection is a good foundation.	No intake, consent, payment or reminder loop.	Complete booking-to-arrival lifecycle.
Start	/start	5.0	Account entry is explicit.	Patient activation depends on unshipped production OTP.	Block external launch until delivery is live.
Login	/login	5.5	Role-aware entry and security states appear considered.	Lint errors sit in this high-risk screen; selection adds cognitive load.	Clean quality gate and reduce first-time ambiguity.
Register organization	/register-org	5.0	Organization model is real.	No clear guided operational setup/migration.	Add stepped setup with readiness confirmation.
Workspace selector	/workspace	6.0	Multi-workspace model is a real differentiator.	Value is niche before customers are multi-clinic.	Make solo path exceptionally simple; reveal complexity progressively.

Screen	Route	Score	Strength	Weakness	Missing information / suggestion
Admin home	/admin	5.3	Owner/admin surface exists.	Admin capability outpaces commercial readiness.	Focus on actionable clinic health, not platform furniture.
Command center	/admin/command-center	5.2	Shows operational attention.	Analytics and resolution actions are thin.	Prioritize outcomes, owner and one-click recovery.

Screen inventory 3 of 6

Screen	Route	Score	Strength	Weakness	Missing information / suggestion
Departments	/admin/departments	5.0	Supports organization structure.	Not a daily buyer value driver.	Keep out of launch demo unless polyclinic fit is proven.
Admin events	/admin/events	4.5	Good engineering observability.	Customer-facing admin value is unclear.	Keep internal unless audit/replay is a paid need.
Admin releases	/admin/releases	3.5	Useful internal release construct.	Does not belong in normal clinic operating narrative.	Protect behind platform-admin role.
Admin settings	/admin/settings	4.8	Basic organization management is necessary.	Plan/clinic nav remains "Soon."	Remove incomplete navigation or finish it.
Admin setup	/admin/setup	5.2	Checklist is directionally correct.	No evidence of activation completion/value metric.	Instrument and make blockers actionable.
Clinic home	/clinic	5.8	Useful consolidated solo surface.	May become a crowded hybrid as features grow.	Define strict solo cockpit information architecture.
Clinic calendar	clinic dashboard	5.0	Calendar foundation exists.	Scheduling depth lacks recurrence/resources.	Do not promise polyclinic operations yet.
Clinic consultation	clinic dashboard	5.5	Shared consult components reduce drift.	Structured clinical model remains immature.	Advance encounter model before expanding templates.
Clinic plan	clinic dashboard	4.0	Plan intent aligns with strategy.	Commercial/treatment planning maturity is incomplete.	Deliver the full plan loop before promotion.
Clinic team	clinic dashboard	5.5	Membership/RBAC foundation is sound.	Nurse/technician daily work is missing.	Align role packaging to real actions.

Screen inventory 4 of 6

Screen	Route	Score	Strength	Weakness	Missing information / suggestion
Staff home	/staff	5.5	Reception surface addresses operational flow.	Navigation can fragment queue/calendar/billing context.	Create a unified front-desk day workspace.
Staff dashboard	/staff/dashboard	3.5	Contains some useful search affordance.	Documented as orphaned/unlinked.	Remove or make it a deliberate landing page.
Queue board	/staff/queue	6.5	Strong stage ownership, wait aging, one-action-per-stage.	No reassign, capacity recovery, global patient search or full calendar context.	Keep as flagship; complete its exception paths.
Walk-in	/staff/walkin	6.0	Minimal capture and identity resolution are valuable.	Family disambiguation and identifier use are weak.	Use Health ID/DOB and explicit duplicate resolution.
Billing desk	/staff/billing	5.7	Checkout, partial/split payment and corrections show real substance.	No online payment, weak search, tax inactive.	Finish patient payment and compliant billing loop.
Staff calendar	/staff/calendar	4.8	Gives a day view.	Day-only, fixed hours, no scheduling power.	Develop before positioning as scheduling solution.
Lab worklist	/staff/lab	4.8	Basic order/result flow exists.	Search, clinical review and integration depth are weak.	Keep narrow or invest in trusted lab workflow.
Doctor Today	/doctor	6.3	"Who is next?" is a strong mental model.	Pause booking is non-functional; no deeper exceptions.	Retain focus and remove false controls.

Screen	Route	Score	Strength	Weakness	Missing information / suggestion
Doctor Workbench	/doctor/workbench	6.4	Three-column context/queue/consult concept is compelling.	Free-text data, autosave confidence and placeholder context limit clinical trust.	Make this the validated clinical flagship.
Doctor Patients	/doctor/patients	4.6	Derived patient list is a start.	Name-only client search and no risk/follow-up facets.	Add secure server search and useful clinical filters.

Screen inventory 5 of 6

Screen	Route	Score	Strength	Weakness	Missing information / suggestion
Doctor schedule	/doctor/schedule	5.0	Day/week/month visibility and blocks exist.	Requests and mature scheduling are absent.	Make scheduling reliable before adding analytics.
Doctor practice	/doctor/practice	3.8	Profile intent is clear.	Many “coming soon” panels make it feel unfinished.	Strip to real booking/profile management.
Doctor profile	/doctor/profile	5.0	Credentials and availability matter.	Awards, memberships and signature are incomplete tiles.	Only show verified, usable attributes.
Patient home	/patient	6.4	Warm, mobile-first framing and next-visit summary are sound.	It cannot proactively communicate or resolve payment.	Turn it into an active care companion.
Patient booking	/patient/book	5.2	Simple discovery and booking path.	Client-side directory search lacks filters and scale.	Build intent-led specialty/location/availability discovery.
Patient doctor profile	/patient/doctors/[id]	5.6	Supports profile-to-slot conversion.	No intake, consent, deposit or expectation setting.	Use pre-visit preparation as conversion advantage.
Patient care	/patient/care	4.5	Care concept aligns with product vision.	The longitudinal care proposition is still thin.	Anchor it in real plans, Rx and communication.
Patient records	/patient/records	5.2	Timeline/Rx/Bills/Tests grouping is understandable.	No record search; depth depends on unstructured visit data.	Ship timeline and metadata search together.
Patient family	/patient/family	5.5	Family profile switching suits India-first use.	Consent/proxy controls are immature.	Add clear guardian and sharing governance.
Patient profile	/patient/profile	5.3	Health summary editing is useful.	Critical safety data quality/review status is unclear.	Add source and last-reviewed signals.

Screen inventory 6 of 6

Screen	Route	Score	Strength	Weakness	Missing information / suggestion
Patient settings	/patient/settings	4.2	Session/privacy intent is appropriate.	Multiple settings remain “coming soon.”	Keep only controls that work.
Patient You	/patient/you	5.0	Clear account framing.	Insurance is shown without supporting workflow.	Remove unimplemented financial claims.
Print document	/print/document/[id]	5.5	Central document-rendering foundation is valuable.	Patient delivery and document breadth are incomplete.	Complete delivery/audit before marketing documents.
Print prescription	/print/prescription/[id]	5.0	Print supports a real clinic behavior.	Rx safety and structured content still lack maturity.	Prioritize signing/amendment and safeguards.
Invite join	/join/[token]	5.4	Provisioned-team onboarding is real.	Role productivity is uneven after acceptance.	Link invitation to role-specific first-day checklist.
Change password	/change-password	5.4	Forced password change is a strong safety control.	No self-service recovery reduces support resilience.	Provide secure reset before scale.

22. User Emotion Journey

Moment	Likely emotion now	Why	Design response
Discover / buy	Curious then skeptical	A strong story meets a broad, unfinished product claim.	Make promise specific and evidence-led.
First setup	Hopeful then uncertain	Organization capability exists; activation/migration path is not proven.	Guide to a first successful clinic day.
Reception rush	Potentially confident	Board makes next state visible.	Add search, reassign, escalation and calendar context.
Doctor consult	Focused then cautious	Workbench is structured visually, but record/safety depth is limited.	Make saving, signing and context trustworthy.
Checkout	Efficient in person	Billing architecture is solid.	Enable patient payment and clear service linkage.
At home patient	Neglected	No external reminder/document/payment/care loop reaches them.	Create a consented after-visit relationship.
Owner review	Underwhelmed	Activity is visible; outcomes/ROI are not.	Show how Auriva improved the month.

23. Customer Delight Opportunities

- A “clinic day closed” moment: every completed encounter shows collected/owed, unresolved exceptions and tomorrow’s risk—not generic dashboard cards.
- One-tap rebook at checkout tied to a plan/follow-up, with patient confirmation sent automatically.
- Patient after-visit card: next action, document, payment status, contact path and reassuring language.
- Reception recovery panel: “Dr. late, 4 waiting >30m, 2 alternative slots” with controlled actions, not an alarm wall.
- Owner value note: “This month Auriva helped recover X follow-ups / Y outstanding balances”—only when data is reliable.
- Structured prescription that prints beautifully, is clearly signed/versioned and lands safely in the patient record.

24. Final Product Report Card

Dimension	Score / 10	Classification
Product	5.2	Red
Design	6.8	Amber
UX	5.4	Amber
Engineering	5.8	Amber
Architecture	6.5	Amber
Scalability	4.2	Red
Business	3.7	Red
Commercial	3.8	Red
Healthcare workflow	5.0	Red
Innovation	5.6	Amber
Overall	5.2	Red

FINAL VERDICT: ⚠ NEEDS SIGNIFICANT WORK Auriva is not ready for commercial launch. It is ready to earn evidence in a very narrow, assisted design-partner pilot only after every P0 recommendation is completed and independently verified.

Appendix C — detailed launch evidence sheets

The following sheets translate the headline audit into testable launch evidence. Each should be owned, rehearsed and signed before a controlled design-partner pilot. They are deliberately operational rather than aspirational.

1. Account creation and clinic setup

TEST INTENT An owner can create a clinic, configure safe defaults, invite the minimum team, add doctors/services and reach a ready-to-operate state without hidden support work.

Evidence observed

The organization and membership model are real. A quantified first-value path, import journey and setup telemetry are not evidenced.

Measurement required

Measure time to ready-to-book, completion/drop-off by setup step, and support contacts per activation.

Pilot exit criterion

A non-technical clinic owner completes a rehearsal with no data ambiguity and receives a clear “ready” confirmation.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

2. Staff and patient identity

TEST INTENT Staff access follows password/change rules; patients authenticate with OTP and select the correct family profile.

Evidence observed

OTP design is thoughtful but production delivery is an open release gate. Identity unification remains future debt.

Measurement required

Record send, delivery, verify, expiry, abuse and failed-login metrics; test account role overlap.

Pilot exit criterion

Live OTP delivery works at target success rate; failures have safe recovery and support traceability.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

3. Online booking and preparation

TEST INTENT A patient finds an appropriate clinician, sees reliable availability, submits required intake/consent, understands price/policy and receives confirmation.

Evidence observed

Slot booking exists. Intake, consent, payment/deposit, delivery and waitlist are absent.

Measurement required

Track conversion by step, abandonment reason, confirmation delivery and no-show rate.

Pilot exit criterion

A pilot patient can book, prepare, cancel/reschedule and be reminded without calling the clinic.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

4. Walk-in identity and registration

TEST INTENT Reception identifies a returning patient or creates a safe minimal profile in seconds.

Evidence observed

Backend supports phone/name/DOB/Health ID, but core UI does not expose Health ID and family disambiguation is weak.

Measurement required

Time identity resolution, duplicate-profile rate and wrong-profile near misses.

Pilot exit criterion

Staff successfully identify family members using two identifiers and can see why a duplicate warning appeared.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

5. Reception queue control

TEST INTENT Reception can see the room state, check in, send in, reassign, recover a wait breach and communicate changes.

Evidence observed

The board clearly shows stage ownership and wait aging. Reassignment, capacity thresholds and outbound communication are missing.

Measurement required

Track wait percentile, queue exceptions, recovery time and manual workaround count.

Pilot exit criterion

A late doctor/absent patient/urgent walk-in scenario is resolved with authorized, audited actions.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

6. Doctor consultation

TEST INTENT A doctor can safely find context, document, prescribe, order tests, save confidently and hand back to checkout.

Evidence observed

The Workbench concept is strong. Clinical record is largely free text; autosave/conflict and clinical governance are insufficiently evidenced.

Measurement required

Measure consult task time, note completion, correction rate and clinician confidence.

Pilot exit criterion

Clinicians complete timed, realistic cases and can identify saved/signed/amended state without ambiguity.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

7. Prescription and diagnostics

TEST INTENT A prescription is structured, signed, versioned, printable/deliverable and safely related to patient history; test results are reviewed and acknowledged.

Evidence observed

Basic Rx and lab flows exist, but medication safety, robust result context and external delivery are absent.

Measurement required

Track Rx completion, amendment, interaction-warning handling and result acknowledgment.

Pilot exit criterion

Clinical Safety Board approves template/content and doctors can demonstrate safe correction/acknowledgment.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

8. Checkout and revenue closure

TEST INTENT Service delivery becomes a transparent invoice, payment is captured in the right channel, a receipt is issued and follow-up is offered.

Evidence observed

Invoice/payment/correction architecture is substantive. Online payment, active tax treatment and service-to-charge integration need completion.

Measurement required

Track completed-to-invoiced, invoiced-to-paid, concession/refund reason and outstanding aging.

Pilot exit criterion

A patient can pay in person or remotely, receive a receipt and be offered the next clinically appropriate booking.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

9. Patient after-visit experience

TEST INTENT Patients receive what matters after care: documents, payment state, next action and a safe contact path.

Evidence observed

The portal has warm navigation and records, but no outbound delivery, online payment or after-visit orchestration.

Measurement required

Track document delivery/open, payment completion, rebooking and care-plan engagement.

Pilot exit criterion

A patient who never reopens the app still receives consented essential information correctly.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

10. Owner decision system

TEST INTENT An owner can see whether the clinic is flowing, collecting and retaining without building spreadsheets.

Evidence observed

Command Center is a foundation; decision-grade KPIs and analytics are thin.

Measurement required

Track wait, no-show, collection, rebooking, plans and staff activity through a single dashboard.

Pilot exit criterion

The owner can answer “what needs attention today?” and “did we improve this month?” in under two minutes.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

11. Multi-role permissions

TEST INTENT Every role sees only safe work, has a usable day-one surface and loses access immediately when changed.

Evidence observed

RBAC/audit foundations are strong; nurse/technician action surfaces and clinic-scoped grants are incomplete.

Measurement required

Run permission matrix regression and role productivity tests.

Pilot exit criterion

No role is sold until its representative can complete the intended workflow end-to-end.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

12. Data integrity and privacy

TEST INTENT Tenant boundaries, audit trails, backups, correction paths and patient sharing protect health data.

Evidence observed

Many safeguards exist; restore evidence, formal privacy/compliance pack, family consent and scale hardening remain gaps.

Measurement required

Test cross-tenant denial, restore RTO/RPO, deletion/retention and audit completeness.

Pilot exit criterion

Security/privacy lead signs a release evidence pack and a restore drill passes.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

13. Accessibility and responsive use

TEST INTENT Keyboard, screen reader, contrast, focus, motion and mobile behavior support real staff and patient diversity.

Evidence observed

Component-level intent exists; independent conformance evidence does not.

Measurement required

Run automated and manual accessibility suite at desktop/tablet/mobile breakpoints.

Pilot exit criterion

Critical journeys complete with keyboard and screen reader; no critical contrast/touch defects remain.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

14. Performance and reliability

TEST INTENT Morning opening, queue spikes, payments, booking and notifications perform predictably under expected load.

Evidence observed

Health/ready endpoints and logging exist; unpaged lists, N+1 reads and single-instance mechanisms are documented risk.

Measurement required

Set SLOs for response time/error rate; load test realistic clinic concurrency.

Pilot exit criterion

SLOs pass with headroom and alert owners are on call for the pilot.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

15. Release quality gates

TEST INTENT The product can be built, tested, deployed, monitored and rolled back reliably.

Evidence observed

Tests/types pass, but lint fails and build relies on remote fonts in this environment.

Measurement required

Make lint/type/test/build, migration, backup, smoke and rollback evidence mandatory.

Pilot exit criterion

A release candidate clears all gates from a clean environment and is rehearsed through rollback.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

16. Customer success and support

TEST INTENT A clinic knows what success means, where to get help and how feedback changes product decisions.

Evidence observed

Release docs have operational notes; formal customer-success model is not evidenced.

Measurement required

Measure adoption, time-to-first-value, support response, workflow completion and NPS/CSAT.

Pilot exit criterion

Each pilot has named sponsor, success plan, escalation channel and weekly value review.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

17. Sales and packaging

TEST INTENT A prospect understands fit, price, implementation, boundaries and proof without a bespoke founder explanation.

Evidence observed

Edition structure is conceptual; price, migration and support offer are not ready.

Measurement required

Track qualified pipeline, conversion, cycle length, implementation time and early retention.

Pilot exit criterion

Sales uses a standard offer for one beachhead, with no capability promises outside shipped scope.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

18. Clinical governance and change control

TEST INTENT Clinical workflow changes have accountable review, risk category, validation and release decision.

Evidence observed

Product correctly defers AI diagnosis; a formal clinical governance operating system is not evidenced.

Measurement required

Audit change records, clinician review, safety incidents and exception approvals.

Pilot exit criterion

No clinical feature ships without documented review, usability validation and post-launch monitoring.

Board challenge

Do not mark this sheet green because a component or API exists. Mark it green only when a representative user can complete the scenario, failure handling is clear, the result is auditable and the operating owner can explain the metric.

Required leadership actions

Product Office Recommendations

- Choose one beachhead and write a one-sentence value proposition with an observable 90-day outcome.
- Move communications and pre-visit readiness ahead of broad feature expansion.
- Use treatment planning only when supported by validated specialty demand.
- Create a decision gate for every milestone: user evidence, success metric, commercial owner and stop condition.

CTO Recommendations

- Restore clean CI: lint/type/test/build; treat production build reproducibility as a release contract.
- Close production configuration, backup/restore, monitoring, shared rate limit and scheduled-job risks.
- Fund encounter data decomposition, search/pagination and multi-clinic correctness before Enterprise claims.
- Publish an evidence-linked operational dashboard and incident runbook.

UX Director Recommendations

- Validate the queue board and Workbench in live clinic simulation before visual polish work.
- Replace every visible false affordance/“Soon” feature with real capability, entitlement-aware hiding or a useful next action.
- Institutionalize keyboard, touch, contrast, screen-reader and responsive acceptance tests.
- Design the entire patient relationship: prepare → attend → pay → receive → follow up.

QA Director Recommendations

- Define P0 patient-safety/privacy/revenue-flow defect categories and stop-ship criteria.
- Automate the full role-based day, including permission denial and failure/retry cases.
- Add visual, accessibility, performance and production synthetic test layers.
- Require a release evidence pack signed by Product, Engineering, Clinical Safety and Operations.

Go-To-Market Recommendations

- Do not sell “Healthcare OS” or Enterprise. Sell an assisted Professional pilot for one clinic type.
- Recruit 3–5 design partners; baseline waiting time, no-shows, collection, rebooking and staff task time.
- Build migration, implementation and support into the offer; do not treat them as afterthoughts.
- Publish only claims that can be demonstrated in the product or measured in pilots.

Launch Checklist

Gate	Exit criterion
P0 product/clinical	OTP, communications, consent strategy, safety governance and E2E paths closed.
P0 engineering	Lint/type/test/build clean; backups restored; alerts verified; production security configuration validated.
P0 commercial	Beachhead, price, contract, onboarding, support owner and pilot KPIs approved.
Pilot readiness	Design partners trained; migration dry-run; sandbox/rehearsal complete; escalation channel staffed.
Launch evidence	One full clinic day completed twice without support intervention; no P0/P1 defects; decision review signed.

Final Launch Decision

Do not authorize commercial launch. Authorize only a controlled, assisted design-partner pilot when P0 gates are green. Reassess after at least three clinics complete 30–90 days of measured usage, with documented evidence of activation, reliability, flow time, collection, rebooking, patient reachability and retention.

Appendix A — current competitor sources

- Cliniko Online Payments: <https://www.cliniko.com/features/transactions/online-payments/>
- Jane Online Booking: <https://jane.app/features/online-booking> ; Jane Features: <https://jane.app/us/features>
- SimplePractice Client Portal: <https://support.simplepractice.com/hc/en-us/articles/207925883> ; billing documents: <https://support.simplepractice.com/hc/en-us/articles/360017844871-How-to-share-billing-documents-with-clients>
- Halaxy Calendar: <https://eu.halaxy.com/feature/calendar>
- Practo Ray: <https://www.practo.com/providers/clinics/ray>
- DocPulse: <https://docpulse.com/>
- Clinicea Features: <https://clinicea.com/features>
- Healthie pricing/positioning: <https://help.gethealthie.com/article/773-healthie-pricing>

Appendix B — audit evidence references

- Release gates and known limitations: docs/RELEASE-CANDIDATE.md; production checklist: docs/production-checklist.md.
- Workflow and UX gaps: docs/product-office-audit/01–05; current maturity: docs/knowledge-base/21-product-maturity-matrix.md.
- Strategic scope and roadmap: docs/AURIVA-PRODUCT-BIBLE.md; docs/AURIVA-RELEASE-1X-ROADMAP.md.
- Risk/debt: docs/technical-debt.md; docs/data-governance.md; prisma/schema.prisma; src/services and src/app/api.